

## Lab Assignment – 2

### Subject: Deep Learning

#### Improving the Neural Network with Fine-Tuning

##### Abstract

A neural network is a network or circuit of neurons, or in a modern sense, an artificial neural network, composed of artificial neurons or nodes. Thus a neural network is either a biological neural network, made up of biological neurons, or an artificial neural network, for solving artificial intelligence problems. For our assignment I used Fashion MINST Dataset with 10 classes. We are fine tuning the model by varying Number of epochs, Number of neurons in the hidden layer, batch size, dropout and optimizers to increase the accuracy for the given dataset.

##### 1. Introduction

A neural network is a network or circuit of neurons, or in a modern sense, an artificial neural network, composed of artificial neurons or nodes. The name has been taken from biology term neurons. A biological neural network is composed of a groups of chemically connected or functionally associated neurons. A single neuron may be connected to many other neurons and the total number of neurons and connections in a network may be extensive. Neural networks are a class of machine learning algorithms used to model complex patterns in datasets using multiple hidden layers and non-linear activation functions. A neural network takes an input, passes it through multiple layers of hidden neurons (mini-functions with unique coefficients that must be learned), and outputs a prediction representing the combined input of all the neurons.

We have used Keras library the assignment. Keras is an open-source software library that provides a Python interface for artificial neural networks. Keras acts as an interface for the TensorFlow library.

##### 2. Methodology

First we loaded the required python libraries required for our assignment. Keras, pandas and PIL are the required ones. We imported the dataset of Fashion MINST which have 60000 training data

and 10000 testing data. We have done total 5 experiments for the assignment. We used sequential model as base model with number of epochs 20. We added two DENSE hidden layers and then we added output layer.

##### 3. Experiments

###### Experiment 1: Normal Model with 2 Hidden Layers and epoch = 20

We used the base model defined above. Here we have used Activation function as “relu” for hidden layers. In relu Activation function we take value as 0 if the output is less than 0, else we take the output value directly. The output layer activation function is “softmax”. Softmax activation function is used to predict a multinomial probability distribution. As we are having 10 output classes we used Softmax as our output activation function. The size of Hidden layers is 128 and output layer is 10. For loss function we used “sparse categorical crossentropy” and optimizer as RMSprop. RMSprop is a gradient based optimization technique used in training neural networks.

Model: "sequential"

| Layer (type)              | Output Shape | Param # |
|---------------------------|--------------|---------|
| dense (Dense)             | (None, 128)  | 100480  |
| activation (Activation)   | (None, 128)  | 0       |
| dense_1 (Dense)           | (None, 128)  | 16512   |
| activation_1 (Activation) | (None, 128)  | 0       |
| dense_2 (Dense)           | (None, 10)   | 1290    |
| activation_2 (Activation) | (None, 10)   | 0       |
| Total params: 118,282     |              |         |
| Trainable params: 118,282 |              |         |
| Non-trainable params: 0   |              |         |

###### Experiment 2: Variation in Epochs (30, 50) With 2 Hidden Layers

Using the same model defined above we changed the number of epochs to 30 and 50.

### No of Epochs = 30

| Model: "sequential_1"     |              |         |
|---------------------------|--------------|---------|
| Layer (type)              | Output Shape | Param # |
| dense_3 (Dense)           | (None, 128)  | 100480  |
| activation_3 (Activation) | (None, 128)  | 0       |
| dense_4 (Dense)           | (None, 128)  | 16512   |
| activation_4 (Activation) | (None, 128)  | 0       |
| dense_5 (Dense)           | (None, 10)   | 1290    |
| activation_5 (Activation) | (None, 10)   | 0       |
| Total params: 118,282     |              |         |
| Trainable params: 118,282 |              |         |
| Non-trainable params: 0   |              |         |

### No of Epochs = 50

| Model: "sequential_2"     |              |         |
|---------------------------|--------------|---------|
| Layer (type)              | Output Shape | Param # |
| dense_6 (Dense)           | (None, 128)  | 100480  |
| activation_6 (Activation) | (None, 128)  | 0       |
| dense_7 (Dense)           | (None, 128)  | 16512   |
| activation_7 (Activation) | (None, 128)  | 0       |
| dense_8 (Dense)           | (None, 10)   | 1290    |
| activation_8 (Activation) | (None, 10)   | 0       |
| Total params: 118,282     |              |         |
| Trainable params: 118,282 |              |         |
| Non-trainable params: 0   |              |         |

### Experiment 3: Changing No of Neurons (64,256) in Each Layer

Using the best model from experiment 2 we changed the Number of neurons in each layer from initial 128 to 64 and 256.

#### N\_Hidden = 64

| Model: "sequential_3"      |              |         |
|----------------------------|--------------|---------|
| Layer (type)               | Output Shape | Param # |
| dense_9 (Dense)            | (None, 64)   | 50240   |
| activation_9 (Activation)  | (None, 64)   | 0       |
| dense_10 (Dense)           | (None, 64)   | 4160    |
| activation_10 (Activation) | (None, 64)   | 0       |
| dense_11 (Dense)           | (None, 10)   | 650     |
| activation_11 (Activation) | (None, 10)   | 0       |
| Total params: 55,050       |              |         |
| Trainable params: 55,050   |              |         |
| Non-trainable params: 0    |              |         |

#### N\_Hidden = 256

| Model: "sequential_4"      |              |         |
|----------------------------|--------------|---------|
| Layer (type)               | Output Shape | Param # |
| dense_12 (Dense)           | (None, 256)  | 200960  |
| activation_12 (Activation) | (None, 256)  | 0       |
| dense_13 (Dense)           | (None, 256)  | 65792   |
| activation_13 (Activation) | (None, 256)  | 0       |
| dense_14 (Dense)           | (None, 10)   | 2570    |
| activation_14 (Activation) | (None, 10)   | 0       |
| Total params: 269,322      |              |         |
| Trainable params: 269,322  |              |         |
| Non-trainable params: 0    |              |         |

### Experiment 4: Changing BatchSize (64,256)

Using the best model from experiment 3 we changed the batch size from initial 128 to 64 and 256.

#### Batch Size = 64

| Model: "sequential_5"      |              |         |
|----------------------------|--------------|---------|
| Layer (type)               | Output Shape | Param # |
| dense_15 (Dense)           | (None, 256)  | 200960  |
| activation_15 (Activation) | (None, 256)  | 0       |
| dense_16 (Dense)           | (None, 256)  | 65792   |
| activation_16 (Activation) | (None, 256)  | 0       |
| dense_17 (Dense)           | (None, 10)   | 2570    |
| activation_17 (Activation) | (None, 10)   | 0       |
| Total params: 269,322      |              |         |
| Trainable params: 269,322  |              |         |
| Non-trainable params: 0    |              |         |

#### Batch Size = 256

| Model: "sequential_6"      |              |         |
|----------------------------|--------------|---------|
| Layer (type)               | Output Shape | Param # |
| dense_18 (Dense)           | (None, 256)  | 200960  |
| activation_18 (Activation) | (None, 256)  | 0       |
| dense_19 (Dense)           | (None, 256)  | 65792   |
| activation_19 (Activation) | (None, 256)  | 0       |
| dense_20 (Dense)           | (None, 10)   | 2570    |
| activation_20 (Activation) | (None, 10)   | 0       |
| Total params: 269,322      |              |         |
| Trainable params: 269,322  |              |         |
| Non-trainable params: 0    |              |         |

## Experiment 5: Improving the net with dropout (0.1, 0.2, 0.3)

Using the best model from experiment 4 we added dropout layer by dropping 10 %, 20 % and 30 % units. Dropout means to dropping out units (both hidden and visible) in a neural network. It refers to ignoring units (i.e. neurons) during the training phase of certain set of neurons which is chosen at random. Dropout is used to prevent a model from overfitting.

### Dropout = 0.1

Model: "sequential\_7"

| Layer (type)               | Output Shape | Param # |
|----------------------------|--------------|---------|
| dense_21 (Dense)           | (None, 256)  | 200960  |
| activation_21 (Activation) | (None, 256)  | 0       |
| dropout (Dropout)          | (None, 256)  | 0       |
| dense_22 (Dense)           | (None, 256)  | 65792   |
| activation_22 (Activation) | (None, 256)  | 0       |
| dropout_1 (Dropout)        | (None, 256)  | 0       |
| dense_23 (Dense)           | (None, 10)   | 2570    |
| activation_23 (Activation) | (None, 10)   | 0       |

=====  
Total params: 269,322  
Trainable params: 269,322  
Non-trainable params: 0

### Dropout = 0.2

Model: "sequential\_8"

| Layer (type)               | Output Shape | Param # |
|----------------------------|--------------|---------|
| dense_24 (Dense)           | (None, 256)  | 200960  |
| activation_24 (Activation) | (None, 256)  | 0       |
| dropout_2 (Dropout)        | (None, 256)  | 0       |
| dense_25 (Dense)           | (None, 256)  | 65792   |
| activation_25 (Activation) | (None, 256)  | 0       |
| dropout_3 (Dropout)        | (None, 256)  | 0       |
| dense_26 (Dense)           | (None, 10)   | 2570    |
| activation_26 (Activation) | (None, 10)   | 0       |

=====  
Total params: 269,322  
Trainable params: 269,322  
Non-trainable params: 0

### Dropout = 0.3

Model: "sequential\_9"

| Layer (type)               | Output Shape | Param # |
|----------------------------|--------------|---------|
| dense_27 (Dense)           | (None, 256)  | 200960  |
| activation_27 (Activation) | (None, 256)  | 0       |
| dropout_4 (Dropout)        | (None, 256)  | 0       |
| dense_28 (Dense)           | (None, 256)  | 65792   |
| activation_28 (Activation) | (None, 256)  | 0       |
| dropout_5 (Dropout)        | (None, 256)  | 0       |
| dense_29 (Dense)           | (None, 10)   | 2570    |
| activation_29 (Activation) | (None, 10)   | 0       |

=====  
Total params: 269,322  
Trainable params: 269,322  
Non-trainable params: 0

## Experiment 5: Testing with Different Optimizers

Here we have used three optimizers RMSprop, Adam, SGD for fitting our model. RMSprop is a gradient based optimization technique used in training neural networks. Adam is a replacement optimization algorithm for stochastic gradient descent for training deep learning models. Stochastic Gradient Descent (SGD) addresses both of these issues by following the negative gradient of the objective after seeing only a single or a few training examples.

### Optimizer = RMSprop

Model: "sequential\_10"

| Layer (type)               | Output Shape | Param # |
|----------------------------|--------------|---------|
| dense_30 (Dense)           | (None, 256)  | 200960  |
| activation_30 (Activation) | (None, 256)  | 0       |
| dropout_6 (Dropout)        | (None, 256)  | 0       |
| dense_31 (Dense)           | (None, 256)  | 65792   |
| activation_31 (Activation) | (None, 256)  | 0       |
| dropout_7 (Dropout)        | (None, 256)  | 0       |
| dense_32 (Dense)           | (None, 10)   | 2570    |
| activation_32 (Activation) | (None, 10)   | 0       |

=====  
Total params: 269,322  
Trainable params: 269,322  
Non-trainable params: 0

### Optimizer = Adam

Model: "sequential\_11"

| Layer (type)               | Output Shape | Param # |
|----------------------------|--------------|---------|
| dense_33 (Dense)           | (None, 256)  | 200960  |
| activation_33 (Activation) | (None, 256)  | 0       |
| dropout_8 (Dropout)        | (None, 256)  | 0       |
| dense_34 (Dense)           | (None, 256)  | 65792   |
| activation_34 (Activation) | (None, 256)  | 0       |
| dropout_9 (Dropout)        | (None, 256)  | 0       |
| dense_35 (Dense)           | (None, 10)   | 2570    |
| activation_35 (Activation) | (None, 10)   | 0       |

=====  
Total params: 269,322  
Trainable params: 269,322  
Non-trainable params: 0

### Optimizer = SGD


Model: "sequential\_12"

| Layer (type)               | Output Shape | Param # |
|----------------------------|--------------|---------|
| dense_36 (Dense)           | (None, 256)  | 200960  |
| activation_36 (Activation) | (None, 256)  | 0       |
| dropout_10 (Dropout)       | (None, 256)  | 0       |
| dense_37 (Dense)           | (None, 256)  | 65792   |
| activation_37 (Activation) | (None, 256)  | 0       |
| dropout_11 (Dropout)       | (None, 256)  | 0       |
| dense_38 (Dense)           | (None, 10)   | 2570    |
| activation_38 (Activation) | (None, 10)   | 0       |

=====  
Total params: 269,322  
Trainable params: 269,322  
Non-trainable params: 0

## 4. Dataset

Fashion-MNIST is a dataset of Zalando's article images consisting of a training set of 60,000 examples and a test set of 10,000 examples. Each example is a 28x28 grayscale image, associated with a label from 10 classes. This image is converted to numpy array of 784 columns.



## 5. Experiments

Different 6 experiments are done to create best accuracy model and the procedure for all six experiments is defined in the methodology section above. The model is trained on the training dataset and then accuracy is checked using testing dataset. Accuracy is checked after every experiment. The results are described in the results section.

| Sr. No. | Model Name       | Train Accuracy | Validation Accuracy | Test Accuracy |
|---------|------------------|----------------|---------------------|---------------|
| 1.      | Simple Model     | 0.8984         | 0.8791              | 0.8881        |
| 2.      | Epoch = 30       | 0.9120         | 0.8811              | 0.8963        |
| 3.      | Epoch = 50       | 0.9296         | 0.8857              | 0.8906        |
| 4.      | N_Hidden = 64    | 0.8906         | 0.8822              | 0.8824        |
| 5.      | N_Hidden = 256   | 0.9301         | 0.9301              | 0.8913        |
| 6.      | Batch Size = 64  | 0.9189         | 0.8802              | 0.8776        |
| 7.      | Batch Size = 256 | 0.9342         | 0.8820              | 0.8899        |
| 8.      | Dropout = 0.1    | 0.9155         | 0.8865              | 0.8909        |

|     |                     |        |        |        |
|-----|---------------------|--------|--------|--------|
| 9.  | Dropout = 0.2       | 0.9043 | 0.8865 | 0.8946 |
| 10. | Dropout = 0.3       | 0.8934 | 0.8854 | 0.8888 |
| 11. | Optimizer = RMSprop | 0.9155 | 0.8867 | 0.8945 |
| 12. | Optimizer = Adam    | 0.9269 | 0.8923 | 0.9031 |
| 13. | Optimizer = SGD     | 0.8609 | 0.8545 | 0.8810 |

## 6. Conclusion

When we increased the no of epochs accuracy increased. Also when we increased Number of Neurons and Batch Size accuracy increased. But when we increased Dropout accuracy Decreased. For RMSprop Optimizer we got 0.8945 accuracy, for Adam Optimizer 0.9031 and for SGD Optimizer we got 0.8810 Accuracy.

## 7. References

- [1] [https://en.wikipedia.org/wiki/Neural\\_network](https://en.wikipedia.org/wiki/Neural_network)
- [2] [https://www.tensorflow.org/datasets/catalog/fashion\\_mnist](https://www.tensorflow.org/datasets/catalog/fashion_mnist)
- [3] <https://www.geeksforgeeks.org/convert-a-numpy-array-to-an-image/>
- [4] <https://machinelearningmastery.com/adam-optimization-algorithm-for-deep-learning/#:~:text=Adam%20is%20a%20replacement%20optimization,sparse%20gradients%20on%20noisy%20problems.>
- [5] <https://machinelearningmastery.com/rectified-linear-activation-function-for-deep-learning-neural-networks/>
- [6] [https://ml-cheatsheet.readthedocs.io/en/latest/nn\\_concepts.html#:~:text=Neural%20networks%20are%20a%20class,and%20non%20](https://ml-cheatsheet.readthedocs.io/en/latest/nn_concepts.html#:~:text=Neural%20networks%20are%20a%20class,and%20non%20)

Dlinear%20activation%20functions.&text=Each%20neuron's%20coefficients%20(weights)%20are,contributed%20to%20the%20total%20error.

[7] <https://machinelearningmastery.com/gradient-descent-with-rmsprop-from-scratch/#:~:text=Root%20Mean%20Squared%20Propagation%2C%20or,step%20size%20for%20each%20parameter.>

[8] <https://machinelearningmastery.com/softmax-activation-function-with-python/#:~:text=The%20softmax%20function%20is%20used%20as%20the%20activation%20function%20in,more%20than%20two%20classes%20labe>